

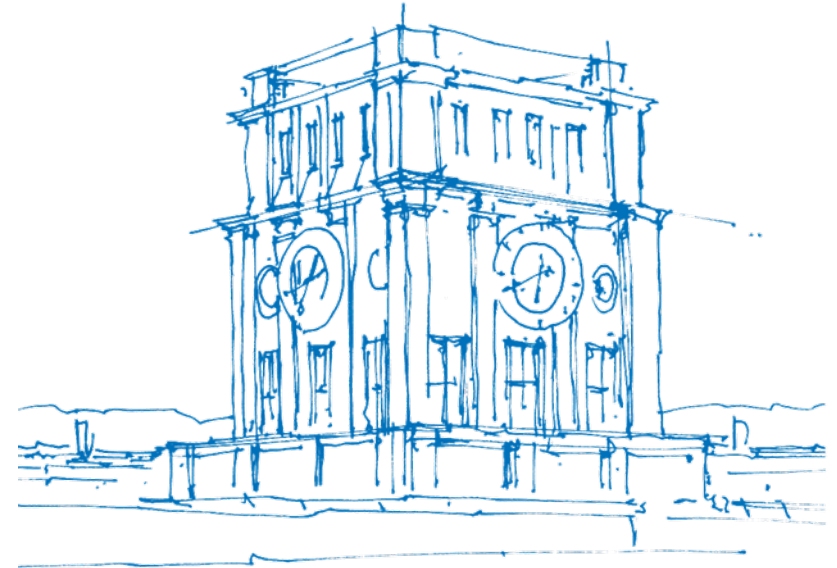
# Parallel Programming Tutorial - Introduction to Pthread API

Amir Raoofy, M.Sc.

Chair for Computer Architecture and Parallel Systems (CAPS)

Technical University Munich

16. April 2018



*TUM Uhrenturm*

# Organization

# Organization

- Course web-page
  - [parprog.lrr.in.tum.de](http://parprog.lrr.in.tum.de)
  - https problem is resolved!
  - Register and login using your @mytum IDs.
  - Course schedule; lecture and tutorial
  - Lecture material and tutorial slides
  - Exercises and assignment submission
- Tutorials: Wednesdays at 8:15 to 9:45
  - Always check the schedule in the web-page
- Where to find us?
  - Chair for Computer Architecture and Parallel Systems (Prof. Dr. M. Schulz)
  - My email address is: [amir.raoofy\(at\)tum.de](mailto:amir.raoofy@tum.de)
  - Room: MI, 01:04:39

Parallel Programming		
Schedule (tentative)		
Week #	Date	Description; lecture and tutorial topics (tentative)
1	09.04.2018	Lecture 01: Introduction / Basic approaches / Types of Parallelism / Metrics / Classification of Languages and Libraries
1	11.04.2018	Lecture 02: Threading (shared memory, general idea, Pthreads, User-level threads)
2	16.04.2018	Tutorial 01: Organization and introduction to Pthreads API
2	18.04.2018	Tutorial 02: More on Pthreads; Synchronization
3	23.04.2018	Lecture 03: OpenMP Basics
3	25.04.2018	Lecture 04: Dependency analysis
4	30.04.2018	Tutorial 03: Parallelism in modern C++
4	02.05.2018	Tutorial 04: Introduction to OpenMP
5	07.05.2018	Lecture 05: OpenMP Advanced topics
5	09.05.2018	Tutorial 05: More on OpenMP; advanced topics
6	14.05.2018	Lecture 06: HPC Architectures / Applications / SuperMUC / Using HPC resources
6	16.05.2018	Tutorial 06: Q/A on OpenMP
7	21.05.2018	Holiday: Pfingstmontag
7	23.05.2018	Tutorial 07: Theoretical; Loop transformations
8	28.05.2018	Lecture 07: MPI Basics
8	30.05.2018	Tutorial 08: Introduction to MPI
9	04.06.2018	Lecture 08: Concurrent Processes Basics / Networking
9	06.06.2018	Lecture 09: MPI Advanced topics
10	11.06.2018	Tutorial 09: More on MPI; advanced topics
10	13.06.2018	Lecture 10: Tuning and Tools
11	18.06.2018	Lecture 11: Scaling / Mapping / ...
11	20.06.2018	Tutorial 10: I.b.a.
12	25.06.2018	Tutorial 11: Q/A on MPI
12	27.06.2018	Tutorial 12: Introduction to Data Centric Models
13	02.07.2018	Lecture 12: Accelerators
13	04.07.2018	Lecture 13: Task based programming and PGAS / Future trends
14	09.07.2018	Tutorial 13: Q/A, exam preparation
14	11.07.2018	Tutorial 14: Q/A, exam preparation

# Assignments

We release two assignments this week and you will have 2 weeks time for completing them!

- We will work on 11 assignments on parallel programming techniques
- Submission of 80% of the assignments brings you 0.3 bonus
- Submission server: <https://parprog.lrr.in.tum.de/Submission/assignments>
  - I will walk you through the submission work-flow at the end of todays tutorial session
- Submissions will be checked for:
  - Plagiarism, correctness (output, threads, synchronization), speedup, memory leaks
- Example solutions will be presented at the following tutorial session
- Topics
  - Pthreads (Posix Threads)
  - C++(11/14/17)
  - OpenMP (Open Multi-Processing)
  - Dependency analysis
  - MPI (Message Passing Interface)

# Assistance on Assignments

Starting this week

- Given by: V.A. Suma and C. Demirsoy, R. Singh
- Emails:
  - vishnu.anilkumar-suma(at)tum.de
  - canberk.demirsoy(at)tum.de
  - rakesh.singh@tum.de
- Room: 01.06.020
- Date and Time:
  - Tuesday 14:00 - 16:00
  - Tuesday 16:00 - 18:00
- If you have questions, write an email to Vishnu and Rakesh and Canberk or visit the assistance sessions

# Resources


- POSIX Threads Programming
- An Introduction to Parallel Programming, by Peter Pacheco
- Programming with Posix Threads, by David Butenhof
- Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill
- Multithreading in Modern C++, by Rainer Grimm
  
- Many thanks to my colleague Andreas Wilhelm
  - Development of the slides and tutorial material

# Course Prerequisites

- knowledge of C/C++ (our code examples and assignments are all in C/C++)
  - memory management
  - pointers /references
  - global vs. static variables
- C/C++(11/14/17) books
  - (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
  - (C99) C Primer Plus, Fifth Edition, by Stephen Prata
  - (C++11/14) The C++ Programming Language, Fourth Edition, by Bjarne Stroustrup
- experience with Linux CLI
  - Book: The Linux Command Line
  - Basic video introduction: The Shell
- knowing compilers/toolchain (e.g., GCC is sufficient here)
  - An Introduction to GCC, by Brian Gough

# Year 2005: The Free Lunch Is Over

- A Fundamental Turn Toward Concurrency in Software
- Software doesn't get (much) faster with the next microprocessor generation
- Developers have to rewrite their software so that multiple computation units are used
- Parallel Programming is hard
  - to write - higher code complexity
  - to do it correctly - easy to introduce bugs
  - to debug - order of thread execution is undefined
  - to make it scalable - will your applications scale with additional cores?
- Qualified developers are necessary → YOU





# Posix Thread Programming

# Posix Thread Programming

## Definition: (Software) Thread

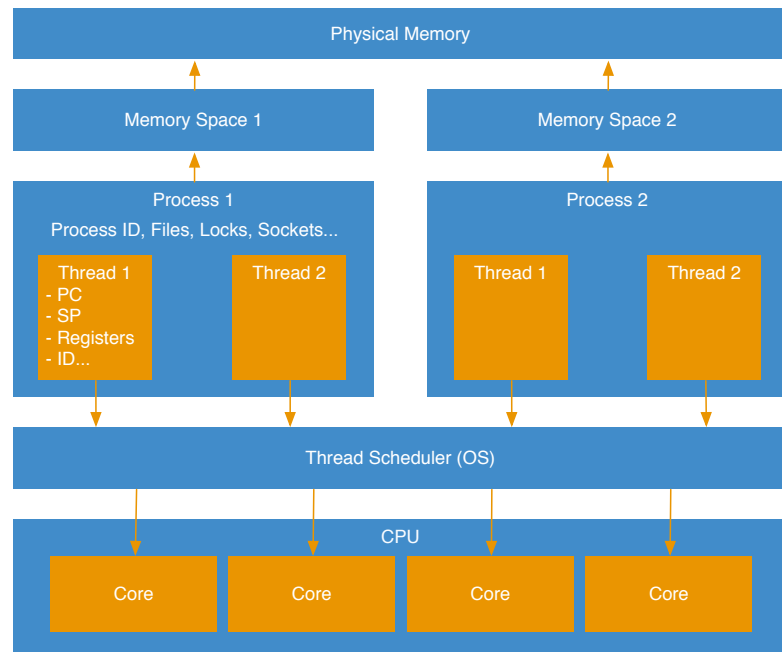
A thread is an independent stream of instructions that can be scheduled to run as such by the operating system. (Own PC and SP)

## POSIX Threads (Pthreads)


- Were defined in 1995 (IEEE Std 1003.1c-1995)
- Is an API that defines a set of types, functions and constants
- Is implemented with a `pthread.h` header and a thread library
- Natively supported by FreeBSD, NetBSD, OpenBSD, Linux, Mac OS X, Android and Solaris
- Functions can be categorized in four groups:
  - Thread management
  - Mutexes
  - Condition variables

# Why use Multithreading?

- **Performance gains**  
Parallel processing by multiple processor cores
- **Increased application throughput**  
Asynchronous system calls possible
- **Increased application responsiveness**  
Application does not need to block operations
- **Replacing process-to-process communications**  
Threads may communicate by shared-memory
- **Efficient use of system resources**  
Lightweight context switches possible
- **Separation of concerns**  
Some problems are inherently concurrent



# Concurrency vs. Parallelism



## Pthread Syntax / Semantics

# Create Pthreads

```
1 int pthread_create(pthread_t *thread ,  
2                   const pthread_attr_t *attr ,  
3                   void *(*start_routine) (void *),  
4                   void *arg );
```

- pthread\_t \*thread,
  - Pointer to thread identifier.
- const pthread\_attr\_t \*attr
  - Optional pointer to pthread\_attr\_t to define behavior, if NULL defaults are used.
- void \*(\*start\_routine) (void \*),
  - Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- void \*arg
  - Pointer to the argument that is used for the executed function.

# Waiting for Pthread to finish

```
1 int pthread_join(pthread_t thread,  
2 void **retval);
```

- pthread\_t thread,
  - Thread identifier, for which this function is waiting.
- void \*\*retval
  - Optional pointer pointing to a void pointer. This can be used to return data of undefined size.

## Example 1; creating a thread

```
1 #include <stdio.h>
2 #include <pthread.h>
3
4 // function to be executed by the thread
5 void* kernel (void* args){
6     printf("hello from the thread!\n");
7     return NULL;
8 }
9
10 int main(int argc, char *argv[])
11 {
12     pthread_t thread;
13     pthread_create(&thread, NULL, kernel, NULL);
14     printf("hello from main\n");
15     pthread_join(thread, NULL);
16
17     return 0;
18 }
```

// allocate a thread  
// create the thread and start executing kernel in parallel  
// wait for the thread to finish executing kernel



# Compile & Output

```
gcc -pthread -Wall -o ex1 ex1.c  
./ex1
```

```
Hello from main!  
Hello from the thread!
```

## Example 2; creating multiple threads

```
1 int main(int argc, char *argv[])
2 {
3     //allocate the threads
4     int num_threads=4;
5     pthread_t *threads = (pthread_t*) malloc (num_threads *sizeof(pthread_t));
6
7     //create threads, start executing kernel in parallel
8     for (int i = 0; i < num_threads; ++i) {
9         pthread_create(&threads[i], NULL, kernel, NULL);
10    }
11
12    //wait for all the threads to finish executing kernel
13    for (int i = 0; i < num_threads; ++i) {
14        pthread_join(threads[i], NULL);
15    }
16
17    free(threads);
18    return 0;
19 }
```

# Output

```
./ex2
```

```
Hello from the thread!  
Hello from the thread!  
Hello from the thread!  
Hello from the thread!
```

## Example 3, passing an argument to threads

```

1 void* kernel (void* args){
2     int id = *(int*)args;
3     printf("Hello from the thread , myid: %d!\n", id);
4     return NULL;
5 }

```

## Example 3, passing an argument to threads (cont.)

```
1  int main(int argc, char *argv[])
2  {
3      //allocate the threads
4      int num_threads=4;
5      pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
6      int* id = (int*) malloc (num_threads*sizeof(int));
7
8      //create threads, start executing kernel in parallel
9      for (int i = 0; i < num_threads; ++i) {
10         id[i]=i; //set the id for the threads
11         pthread_create(&threads[i], NULL, kernel, id+i); //pass the id as argument to the threads
12     }
13
14     //wait for all the threads to finish executing kernel
15     for (int i = 0; i < num_threads; ++i) {
16         pthread_join(threads[i], NULL);
17     }
18
19     free(threads); free(id);
20     return 0;
21 }
```

# Output

```
./ex3
```

```
Hello from the thread, myid: 1!  
Hello from the thread, myid: 0!  
Hello from the thread, myid: 2!  
Hello from the thread, myid: 3!
```

## Example 4; process and thread IDs

```
1 void* kernel (void* args){  
2     int id = *(int*)args;  
3     printf("Hello from the thread , myid: %d, PID: %d, TID:%d!\n", id , getpid(), (int) gettid());  
4     return NULL;  
5 }
```

# Output

```
./ex4
```

```
Hello from the thread, myid: 1, PID: 12347, TID:12349!  
Hello from the thread, myid: 0, PID: 12347, TID:12348!  
Hello from the thread, myid: 2, PID: 12347, TID:12350!  
Hello from the thread, myid: 3, PID: 12347, TID:12351!
```



## Example 5, passing multiple arguments

```
1 struct pthread_args
2 {
3     long thread_id ;
4     long num_threads ;
5 };
6
7 void* kernel (void* args){
8     struct pthread_args *arg = (struct pthread_args*) args;
9     printf("Hello from the thread, number of threads: %ld, myid: %ld, PID: %d, TID:%d!\n", \
10         arg->num_threads, arg->thread_id, getpid(), (int) gettid());
11     return NULL;
12 }
```

## Example 5, passing multiple arguments (cont.)

```
1 int main(int argc, char *argv[])
2 {
3     int num_threads=4;
4     pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5     struct pthread_args* args = (struct pthread_args*) malloc (num_threads*sizeof (struct pthread_args));
6
7     for (int i = 0; i < num_threads; ++i) {
8         //set the id and num threads in args for the threads
9         args[i].thread_id=i;
10        args[i].num_threads=num_threads;
11        //pass the args as argument to the threads
12        pthread_create(&threads[i], NULL, kernel, args+i); // passing args[i] to threads[i]
13    }
14
15    for (int i = 0; i < num_threads; ++i) {
16        pthread_join(threads[i], NULL);
17    }
18
19    free(threads); free(args);
20    return 0;
21 }
```

## Example 6, how to get data out of threads

```
1 struct pthread_args
2 {
3     int in ;
4     int out ;
5 };
6
7 void* kernel_double (void* args){
8     struct pthread_args *arg = (struct pthread_args*) args;
9     arg->out = 2*arg->in;
10    return NULL;
11 }
```

## Example 6, how to get data out of threads (cont.)

```
1 int main(int argc, char *argv[])
2 {
3     int num_threads=4;
4     pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5     struct pthread_args* args = (struct pthread_args*) malloc (num_threads*sizeof (struct pthread_args));
6
7     for (int i = 0; i < num_threads; ++i) {
8         args[i].in=i; //set the input in args
9         pthread_create(&threads[i], NULL, kernel_double, args+i);
10    }
11
12    for (int i = 0; i < num_threads; ++i) {
13        pthread_join(threads[i], NULL);
14    }
15
16    for (int i = 0; i < num_threads; ++i) {
17        printf("Double of %d is %d!\n", args[i].in, args[i].out);
18    }
19
20    free (threads); free (args); return 0;
21 }
```

## Example 7, return data from threads

```
1 void* kernel_double (void* args){  
2     int in = *(int*) args;  
3     int *out = (int*) malloc (1*sizeof (int));  
4     *out = 2*in;  
5     return (void*)out;  
6 }
```

## Example 7, return data from threads (cont.)

```
1 int main(int argc, char *argv[])
2 {
3     int num_threads=4;
4     pthread_t *threads = (pthread_t*) malloc (num_threads*sizeof(pthread_t));
5     int* in = (int*) malloc (num_threads*sizeof(int));
6
7     for (int i = 0; i < num_threads; ++i) {
8         in[i]=i; //set the input for the threads
9         pthread_create(&threads[i], NULL, kernel_double, in+i);
10    }
11
12    for (int i = 0; i < num_threads; ++i) {
13        int *out;
14        pthread_join(threads[i], (void*)&out);
15        printf("Double of %d is %d!\n", in[i], *out);
16        free(out);
17    }
18
19    free (threads); free (in); return 0;
20 }
```

# What have we covered so far?

- Creating new threads with `pthread_create`
- Waiting for threads to finish with `pthread_join`
- Passing arguments to a pthread function
- Returning results from pthread function

## Assignment 1: “Mandelbrot set” in parallel



# Assignment: Mandelbrot


Starting this week, you have two weeks time.

- Use Pthreads to parallelize mandelbrot\_draw()
- Your solution should have a speedup greater than 3.0 using 4 threads

```

1 void mandelbrot_draw( ... some args ) {
2     ...
3     for (int i = 0; i < y_resolution; i++)
4     {
5         for (int j = 0; j < x_resolution; j++)
6         {
7             //embarrassingly parrallel calculation of pixels
8             ...
9         }
10    }
11 }

```



# Assignment: Mandelbrot (cont.)

## Build the program

- Makefile:  
make

## Usage of the program

- Sequential:  
`./mandelbrot_set_seq -h`
- Parallel:  
`./mandelbrot_set_par -t 4 -r 480x380 -i 1000 -v [-2.0,0.5]x[-1.25,1.25] -f mandelbrot.ppm`

# Assignment: Mandelbrot - provided files

- Makefile
  - contains rules to build executables
  - available targets: parallel, sequential, all (default), clean
  - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
  - main function - argument handling + file handling + call draw\_mandelbrot()
- mandelbrot\_set.h
  - Header file for mandelbrot\_set\_\*.c
- mandelbrot.c
  - Defines helper functions
- mandelbrot\_set\_seq.c
  - Sequential version of draw\_mandelbrot()
- student/mandelbrot\_set\_par.c
  - Implement the parallel version in this file

# Assignment: Mandelbrot - provided files (cont.)

- `unit_test.c`
  - The unit tests that execute both the serial and parallel version to compare results.

# Assignment: Extract, Build, and Run

1. Extract all files to the current directory  
`tar -xvf assignment1.tar.gz`
2. Build the program  
`make [sequential] [parallel] [unit_test]`
  - sequential: build the sequential program
  - parallel: build the parallel program
  - unit\_test: builds the unit tests
3. Run the sequential program (with default parameters)  
`student/mandelbrot_set_seq`
4. Run the parallel program (with 4 threads)  
`student/mandelbrot_set_par -t 4`

# Submission

1. Log into the website
2. Go to Assignments
3. Use the link for Assignment 1
4. Upload your `mandelbrot_set_par.c` file
5. Press Submit

Parallel Programming
Home
Lecture material
Tutorial slides
Assignments
Schedule
Contacts
Impressum
Welcome ga59how Logout

Browse... No file selected.
Submit

Assignment Name	Submitted By	Date Of Submission	Status
Pthread 1; parallel Mandelbrot set	Amir Raoofy	April 16, 2018, 1 a.m.	✓ Accepted

Build step successful!  
Correctness checks successful!  
Pthread checks successful!  
Runtime checks successful! - speedup: 3.511  
Parallel Runtime: 0.88248

**Top 10 Runtimes**  
0.88248